

End of preview. [Expand](#) in  Dataset Viewer.

< Previous **1** 2 3 ... 1,140 Next >

Open Thoughts

 Made with Curator

Open-Thoughts-114k

Open synthetic reasoning dataset with 114k high-quality examples covering math, science, code, and puzzles!

Available Subsets


default subset containing ready-to-train data used to finetune the [OpenThinker-7B](#) and [OpenThinker-32B](#) models:

```
ds = load_dataset("open-thoughts/OpenThoughts-114k",
```

metadata subset containing extra columns used in dataset construction:

- problem
- ground_truth_solution
- deepseek_reasoning
- deepseek_solution
- domain
- source
- test_cases (code only)
- starter_code (code only)


```
ds = load_dataset("open-thoughts/OpenThoughts-114k",
```

 bartowski/OpenThinker-...


 Text Generation • • ⬇ 2.07k • <

 mradermacher/OpenThink...

Updated 5 days ago • ⬇ 733 • ♥ 1

 mradermacher/Reasoning...

Updated 8 days ago • ⬇ 721 • ♥ 1

 mradermacher/Phi4.Turn...

Updated 12 days ... • ⬇ 522 • ♥ 1

[Browse 54 models trained on this dataset](#)

OpenThinker Models

The numbers reported in the tables below are evaluated with our open-source tool [Evalchemy](#).

	AIME24	MATH500	GPQA-Diamond	LCBv2 Easy	LCBv2 Medium	LCBv2 Hard
OpenThinker-32B	66	90.6	61.6	95.1	70.9	26.6
OpenThinker-7B	31.3	83.0	42.4	75.3	28.6	6.1
Bespoke-Stratos-7B	22.7	79.6	38.9	71.4	25.2	0.0
DeepSeek-R1-Distill-Qwen-7B	60	88.2	46.9	79.7	45.1	14.0
gpt-4o-0513	8.7	75.8	46.5	87.4	42.7	8.0
o1-mini	64	85.6	60	92.8	74.7	39.0

We are fully open-source. Our [model weights](#), [datasets](#), [data generation code](#), [evaluation code](#), and [training code](#) are all publicly available.

	Open Weights	Open Data	Open Code
OpenThinker-32B	✓	✓	✓
OpenThinker-7B	✓	✓	✓
Bespoke-Stratos-7B	✓	✓	✓
DeepSeek-R1-Distill models	✓	✗	✗

	Open Weights	Open Data	Open Code
OpenAI/Gemini	✗	✗	✗

We are actively working towards improving the dataset, so please stay tuned!

Data Curation Recipe

Code

- [BAAI/TACO](#)
- [codeparrot/apps](#)
- [deepmind/code_contests](#)
- [MatrixStudio/Codeforces-Python-Submissions](#)

Math

- [AI-MO/NuminaMath-CoT](#)

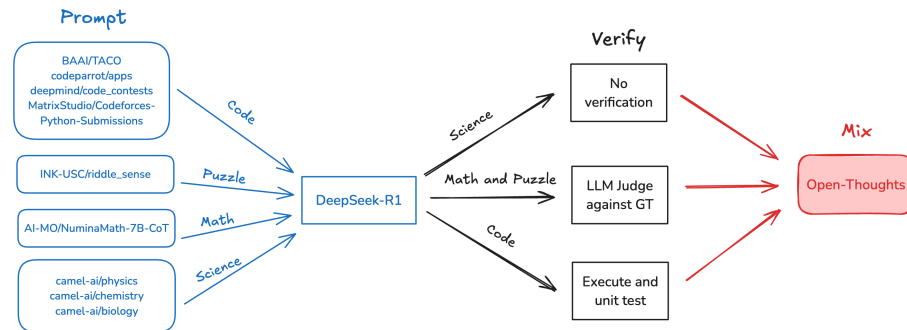
Science

- [camel-ai/chemistry](#)
- [camel-ai/biology](#)
- [camel-ai/physics](#)

Puzzle

- [INK-USC/riddle_sense](#)

Using a curated mix of the datasets above, we generate reasoning traces from DeepSeek-R1 and verify correctness to construct the final dataset.



The full code for the data generation pipeline is publicly available [in our github repo](#).

Citation

```
@misc{openthoughts,  
  author = {Team, OpenThoughts},  
  month = jan,  
  title = {{Open Thoughts}},  
  howpublished = {https://open-thoughts.ai},  
  year = {2025}  
}
```

Links

- 🇮🇹 [OpenThinker-32B Blog Post](#)
- 🇮🇹 [Measuring Reasoning with Evalchemy Blog Post](#)
- 🇮🇹 [Open Thoughts Launch Blog Post](#)
- 🖥️ [Open Thoughts GitHub Repository](#)
- 🧠 [OpenThoughts-114k dataset](#) - this dataset.
- 🤖 [OpenThinker-32B model](#)
- 🤖 [OpenThinker-7B model](#)
- 🇮🇹 [Bespoke-Stratos Blog Post](#)
- 🧠 [Bespoke-Stratos-17k dataset](#)
- 🤖 [Bespoke-Stratos-32B model](#)

System theme

[TOS](#) [Privacy](#) [Stratos-7B model](#)

[About](#)

[Jobs](#)



[Models](#)

[Datasets](#)

[Spaces](#)

[Pricing](#)

[Docs](#)

Visualization

All 114k examples, clustered by semantic similarity, can be explored in [Nomic Atlas](#).

